



Week 10 Lecture 1

▼ Class	BSCCS2001
🕒 Created	@November 8, 2021 1:22 PM
🔗 Materials	
≡ Module #	46
▼ Type	Lecture
# Week #	10

Transactions

Transaction Concept

- A transaction is a unit of program execution that accesses and possibly updates various data items
- For example: transaction to transfer \$50 from account A to account B
 - read(A)
 - $A := A - 50$
 - write(A)
 - read(B)
 - $B := B + 50$
 - write(B)
- Two main issue to deal with:
 - Failures of various kinds, such as hardware failure and system crash
 - Concurrent execution of multiple transactions

Required properties of a Transaction: ACID: Atomicity

- **Atomicity Requirement**
 - If the transaction fails after step 3 and before step 6, money will be "lost" leading to an inconsistent database state
 - Failure could be due to software or hardware
 - The system should ensure that updates of a partially executed transaction are not reflected in the database

Required properties of a Transaction: ACID: Consistency

- Consistency Requirement
 - $A + B$ must be unchanged by the execution of the transaction
 - In general, consistency requirements include
 - Explicitly specified integrity constraints
 - primary keys and foreign keys
 - Implicit integrity constraints
 - sum of balances of all accounts, minus sum of loan amounts must equal value of cash-in-hand
 - A transaction, when starting to execute, must see a consistent database
 - During transaction execution the database may be temporarily inconsistent
 - When the transaction completes successfully the database must be consistent
 - Erroneous transaction logic can lead to inconsistent

Required properties of a Transaction: ACID: Isolation

- Isolation Requirement
 - If between steps 3 and 6 (of the fund transfer transaction), another transaction T2 is allowed to access the partially updated database, it will see an inconsistent database (the sum $A + B$ will be less than it should be)

T1	T2
1. read(A)	
2. $A := A - 50$	
3. write(A)	
	read(A), read(B), print(A + B)
4. read(B)	
5. $B := B + 50$	
6. write(B)	

- Isolation can be ensured trivially by running transactions serially
 - That is, one after the other
- However, executing multiple transactions concurrently has significant benefits

Required properties of a Transaction: ACID: Durability

- Durability Requirement
 - Once the user has been notified that the transaction has completed (that is, the transfer of \$50 has taken place), the updates to the database by the transaction must persist even if there are software or hardware failures

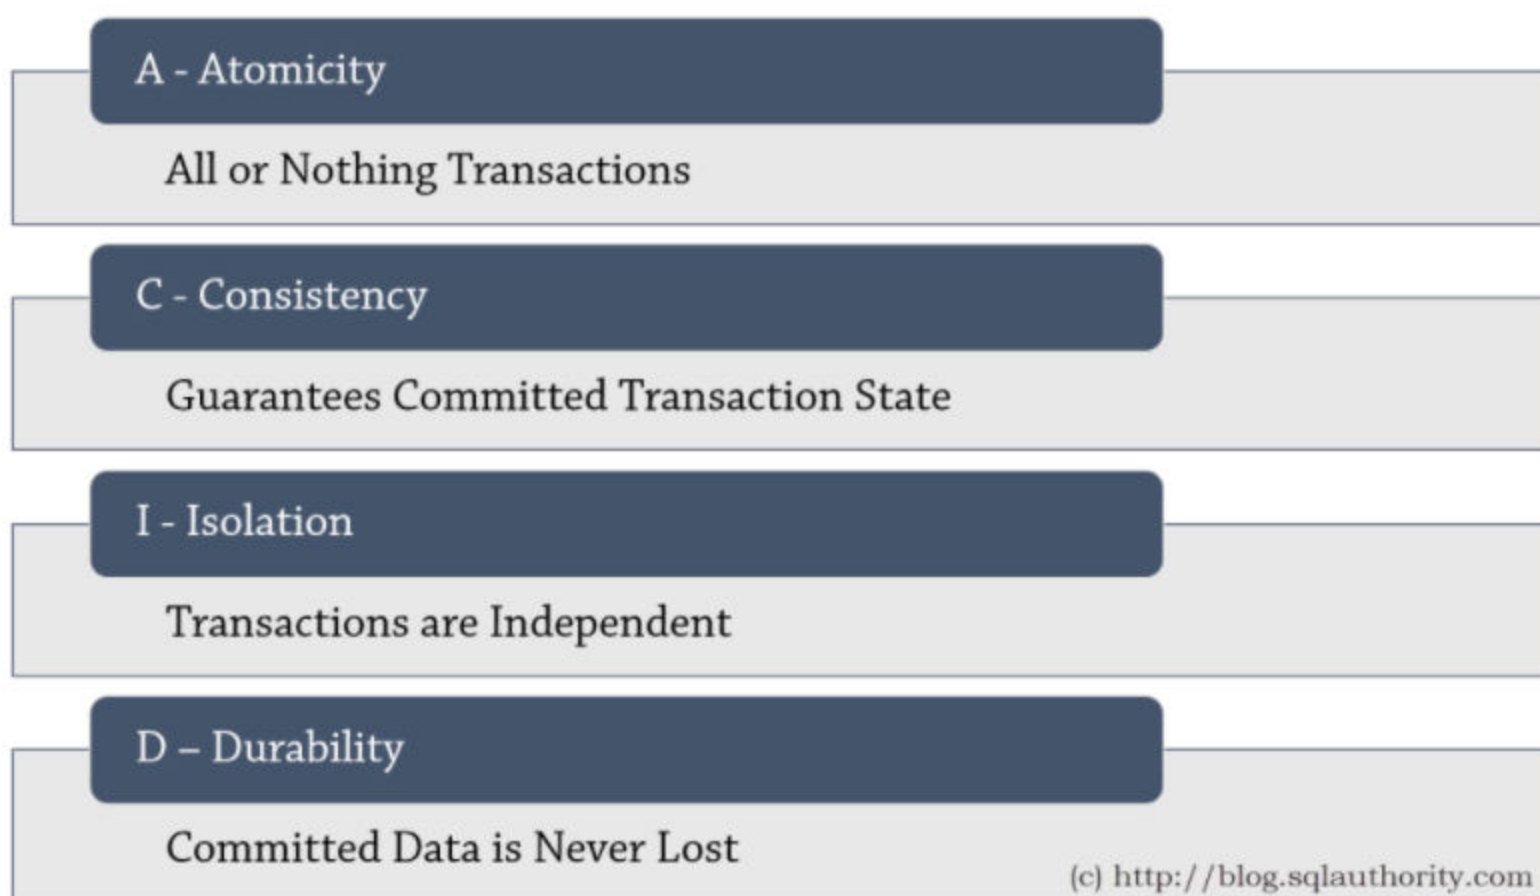
ACID Properties

A **transaction** is a unit of program execution that accesses and possibly updates various data items

- **Atomicity:** Atomicity guarantees that each transaction is treated as a single unit, which either succeeds completely or fails completely
 - If any of the statements constituting a transactions fails to complete, the entire transaction fails and the database is left unchanged
 - Atomicity must be guaranteed in every situation, including power failures, errors and crashes
- **Consistency:** Consistency ensures that a transaction can only bring the database from one valid state to another, maintaining database invariants
 - Any data written to the database must be valid according to all defined rules, including constraints, cascades, triggers and any combination thereof

- **Isolation:** Transactions are often executed concurrently (multiple transactions reading and writing to a table at the same time)
 - Isolation ensures that concurrent execution of transactions leaves the database in the same state that would have been obtained if the transactions were executed sequentially
- **Durability:** Durability guarantees that once a transactions has been committed, it will remain committed even in the case of a system failure (like power outage or crash)
 - This usually means that completed transactions (or their effects) are recorded in non-volatile memory

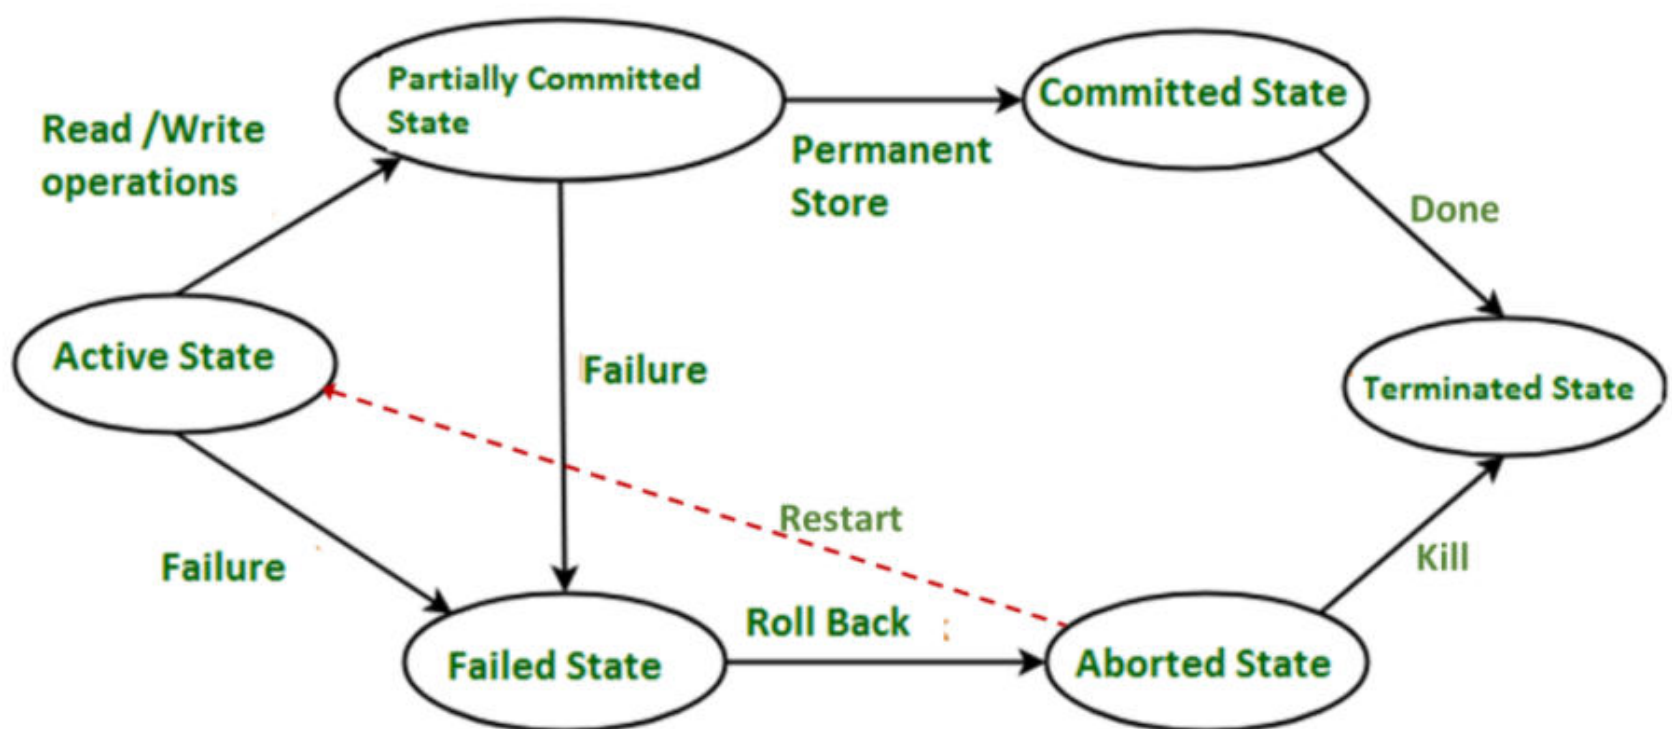
ACID Properties: Quick Reckoner



Transaction States

- Every transaction can be in one of the following states (like Process States in OS)
 - **Active**
 - The initial state; the transaction stays in this state while it is executing
 - **Partially committed**
 - After the final statement has been executed
 - **Failed**
 - After the discovery that normal execution can no longer proceed
 - **Aborted**
 - After the transaction has been rolled back and the database restored to its state prior to the start of the transaction
 - Two options after it has been aborted
 - Restart the transaction: Can be done only if no internal logical error
 - Kill the transaction
 - **Committed**
 - After successful completion
 - **Terminated**
 - After it has been committed or aborted (killed)

Transitions for Transaction states



Concurrent Executions

- Multiple transactions are allowed to run concurrently in the system
 - Advantages are:
 - Increased processor and disk utilization**, leading to better transaction throughput
 - For example, one transaction can be using the CPU while another is reading from or writing to the disk
 - Reduced average response time** for transactions: short transactions need not wait behind long ones
- Concurrency Control Schemes:** Mechanisms to achieve isolation
 - To control the interaction among the concurrent transactions in order to prevent them from destroying the consistency of the database

Schedules

- Schedules:** A sequence of instructions that specify the chronological order in which instructions of concurrent transactions are executed
 - A schedule for a set of transactions must consist of all instructions of those transactions
 - Must preserve the order in which the instructions appear in each individual transaction
- A transaction that successfully completes its execution will have a commit instruction as the last statement
 - By default, transaction assumed to execute commit instruction as its last step
- A transaction that fails to successfully complete its execution will have an abort instruction as the last statement

Schedule 1

- Let T_1 transfer \$50 from A to B and T_2 transfer 10% of the balance from A to B
- An example of a serial schedule in which T_1 is followed by T_2

T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit

A	B	A+B	Transaction	Remarks
100	200	300	@ Start	
50	200	250	T1, write A	
50	250	300	T1, write B	@ Commit
45	250	295	T2, write A	
45	255	300	T2, write B	@Commit

Consistent @ Commit
Inconsistent @ Transit
Inconsistent @ Commit

Schedule 2

- A serial schedule in which T_2 is followed by T_1

T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit

A	B	A+B	Transaction	Remarks
100	200	300	@ Start	
90	200	290	T2, write A	
90	210	300	T2, write B	@ Commit
40	210	250	T1, write A	
40	260	300	T1, write B	@Commit

Consistent @ Commit
Inconsistent @ Transit
Inconsistent @ Commit

Values of A & B are different from
Schedule 1 – yet consistent

Schedule 3

- Let T_1 and T_2 be the transactions defined previously
- The following schedule is not a serial schedule, but it is equivalent to Schedule 1

Schedule 3		Schedule 1	
T_1	T_2	T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	 read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit	read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	 read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit

A	B	A+B	Transaction	Remarks
100	200	300	@ Start	
50	200	250	T1, write A	
45	200	245	T2, write A	
45	250	295	T1, write B	@ Commit
45	255	300	T2, write B	@Commit

- Consistent @ Commit
- Inconsistent @ Transit
- Inconsistent @ Commit

Note – In schedules 1, 2 and 3, the sum " $A + B$ " is preserved

Schedule 4

- The following concurrent schedule does not preserve the sum of " $A + B$ "

T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	 read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit

A	B	A+B	Transaction	Remarks
100	200	300	@ Start	
90	200	290	T2, write A	
50	200	250	T1, write A	
50	250	300	T1, write B	@ Commit
50	210	260	T2, write B	@ Commit

- Consistent @ Commit
- Inconsistent @ Transit
- Inconsistent @ Commit



Week 10 Lecture 2

▼ Class	BSCCS2001
🕒 Created	@November 8, 2021 1:55 PM
🔗 Materials	
☰ Module #	47
▼ Type	Lecture
# Week #	10

Transactions: Serializability

Serializability

- **Assumption:** *Each transaction preserves database consistency*
- Thus, serial execution of a set of transactions preserves database consistency
- A (possible concurrent) schedule is serializable if it is equivalent to a serial schedule
- Different forms of schedule equivalence give rise to the notions of:
 - Conflict Serializability
 - View Serializability

Recap Schedule 3: Serializable

- Let T_1 and T_2 be the transactions defined previously
- The following schedule is not a serial schedule, but it is equivalent to Schedule 1

Schedule 3		Schedule 1						
T_1	T_2	T_1	T_2	A	B	A+B	Transaction	Remarks
read (A) $A := A - 50$ write (A)		read (A) $A := A - 50$ write (A)		100	200	300	@ Start	
	read (A) $temp := A * 0.1$ $A := A - temp$ write (A)	read (B) $B := B + 50$ write (B) commit		50	200	250	T1, write A	
				45	200	245	T2, write A	
read (B) $B := B + 50$ write (B) commit			read (A) $temp := A * 0.1$ $A := A - temp$ write (A)	45	250	295	T1, write B	@ Commit
	read (B) $B := B + temp$ write (B) commit		read (B) $B := B + temp$ write (B) commit	45	255	300	T2, write B	@Commit

Consistent @ Commit

Inconsistent @ Transit

Inconsistent @ Commit

Note: In schedules 1, 2 and 3, the sum "A + B" is preserved

Recap Schedule 4: Not Serializable

- The following concurrent schedule does not preserve the sum of "A + B"

T_1	T_2					
read (A) $A := A - 50$		A	B	A+B	Transaction	Remarks
	read (A) $temp := A * 0.1$ $A := A - temp$ write (A)	100	200	300	@ Start	
	read (B)	90	200	290	T2, write A	
write (A) read (B) $B := B + 50$ write (B) commit		50	200	250	T1, write A	
	$B := B + temp$ write (B) commit	50	250	300	T1, write B	@ Commit
		50	210	260	T2, write B	@ Commit

Consistent @ Commit

Inconsistent @ Transit

Inconsistent @ Commit

Simplified View of Transactions

- We ignore operations other than read and write instructions
 - Other operations happen in memory (are temporary in nature) and (mostly) do not affect the state of the database
 - This is a simplifying assumption for analysis
- We assume that transactions may perform arbitrary computations on data in local buffers in between reads and writes
- Our simplified schedules consist of only read and write instructions

Conflicting Instructions

- Let I_i and I_j be 2 instructions from transactions T_i and T_j respectively
- Instructions I_i and I_j conflict if and only if there exists some item Q accessed by both I_i and I_j and at least one of these instructions write to Q
 - $I_i = \text{read}(Q), I_j = \text{read}(Q) \rightarrow I_i$ and I_j don't conflict
 - $I_i = \text{read}(Q), I_i = \text{write}(Q) \rightarrow$ They conflict
 - $I_i = \text{write}(Q), I_j = \text{read}(Q) \rightarrow$ They conflict
 - $I_i = \text{write}(Q), I_j = \text{write}(Q) \rightarrow$ They conflict
- Intuitively, a conflict between I_i and I_j forces a (logical) temporal order between them

- If I_i and I_j are consecutive in a schedule and they do not conflict, their results would remain the same even if they had been interchanged in the schedule

Conflict Serializability

- If a schedule S can be transformed into a schedule S' by a series of swaps of non-conflicting instructions, we say that S and S' are conflict equivalent
 - We say that a schedule S is conflict serializable if it is conflict equivalent to a serial schedule
-
- Schedule 3 can be transformed into Schedule 6, a serial schedule where T_2 follows T_1 by a series of swaps of non-conflicting instructions:
 - Swap $T1.read(B)$ and $T2.write(A)$
 - Swap $T1.read(B)$ and $T2.read(A)$
 - Swap $T1.write(B)$ and $T2.write(A)$
 - Swap $T1.write(B)$ and $T2.read(A)$

These swaps do not conflict as they work with different items (A or B) in different transactions

T_1	T_2
read (A) write (A)	
	read (A) write (A)
read (B) write (B)	
	read (B) write (B)

Schedule 3

T_1	T_2
read(A) write(A)	
read(B)	read(A)
write(B)	write(A)
	read(B) write(B)

Schedule 5

T_1	T_2
read (A) write (A) read (B) write (B)	
	read (A) write (A) read (B) write (B)

Schedule 6

- Example of a schedule that is not conflict serializable

T_3	T_4
read (Q)	
write (Q)	write (Q)

- We are unable to swap instructions in the above schedule to obtain either the serial schedule $\langle T_3, T_4 \rangle$ or the serial schedule $\langle T_4, T_3 \rangle$

Example: Bad Schedule

Consider two transactions:

Transaction 1

UPDATE accounts

SET balance = balance - 100

WHERE acct_id = 31414

Transaction 2

UPDATE accounts

SET balance = balance * 1.005

	A	B
(initial:)	200.00	100.00
$r_1(A)$:		
$r_2(A)$:		
$w_1(A)$:	100.00	
$w_2(A)$:	201.00	
$r_2(B)$:		
$w_2(B)$:		100.50

Schedule S

- In terms of read/write, ~~we have no read/write~~, we can write this as:

Transaction 1: $r_1(A), w_1(A) // A$ is the balance for $acct_id = 31414$

Transaction 2: $r_2(A), w_2(A), r_2(B), w_2(B) // B$ is the balance of other accounts

- Consider schedule S:
 - Schedule S: $r_1(A), r_2(A), w_1(A), w_2(A), r_2(B), w_2(B)$
 - Suppose: A starts with \$200 and account B starts with \$100
- Schedule S is very bad!
 - We withdrew \$100 from account A, but somehow the database has recorded that our account now holds \$201

- Ideal schedule is serial:

Serial schedule 1:

 $r_1(A), w_1(A), r_2(A), w_2(A), r_2(B), w_2(B)$

Serial schedule 2:

 $r_2(A), w_2(A), r_2(B), w_2(B), r_1(A), w_1(A)$
- We call a schedule **serializable** if it has the same effect as some serial schedule regardless of the specific information in the database
- As an example, consider Schedule T, which has swapped the third and fourth operations from S:
 - Schedule S: $r_1(A), r_2(A), w_1(A), w_2(A), r_2(B), w_2(B)$
 - Schedule T: $r_1(A), r_2(A), w_2(A), w_1(A), r_2(B), w_2(B)$
- By first example, the outcome is the same as Serial schedule 1
 - But that's just a peculiarity of the data, as revealed by the second example, where the final value of A can't be the consequence of either of the possible serial schedules
- So, neither S nor T are serializable

T1	Schedule 1: T1-T2		Schedule 2: T2-T1	
T2	A	B	A	B
Initial Value	200.00	100.00	200.00	100.00
Final Value	100.00	100.00	201.00	100.50
Initial Value	100.00	100.00	201.00	100.50
Final Value	100.50	100.50	101.00	100.50

A is \$100 initially			A is \$200 initially		
	A	B		A	B
(initial:)	100.00	100.00	(initial:)	200.00	100.00
$r_1(A)$:			$r_1(A)$:		
$r_2(A)$:			$r_2(A)$:		
$w_2(A)$:	100.50		$w_2(A)$:	201.00	
$w_1(A)$:	0.00		$w_1(A)$:	100.00	
$r_2(B)$:			$r_2(B)$:		
$w_2(B)$:		100.50	$w_2(B)$:		100.50

Schedule T

Example: Good Schedule

- What's a non-serial example of serializable schedule?
 - We could credit interest to A first then withdraw the money, then credit interest to B:
 - Schedule U: $r_2(A), w_2(A), r_1(A), w_1(A), r_2(B), w_2(B)$
 - Initial: A = 200, B = 100
 - Final: A = 101, B = 100.50

- Schedule U is conflict serializable to Schedule 2:

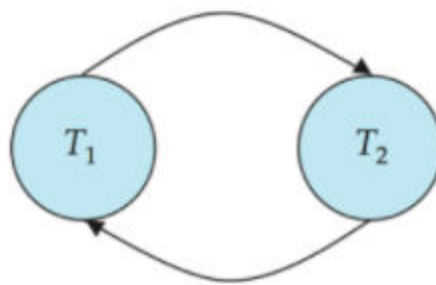
Schedule U : $r_2(A), w_2(A), r_1(A), w_1(A), r_2(B), w_2(B)$
 swap $w_1(A)$ and $r_2(B)$: $r_2(A), w_2(A), r_1(A), r_2(B), w_1(A), w_2(B)$
 swap $w_1(A)$ and $w_2(B)$: $r_2(A), w_2(A), r_1(A), r_2(B), w_2(B), w_1(A)$
 swap $r_1(A)$ and $r_2(B)$: $r_2(A), w_2(A), r_2(B), r_1(A), w_2(B), w_1(A)$
 swap $r_1(A)$ and $w_2(B)$: $r_2(A), w_2(A), r_2(B), w_2(B), r_1(A), w_1(A)$: Schedule 2

Serializability

- Are all serializable schedules conflict-serializable? No
- Consider the following schedule for a set of three transactions
 - $w_1(A), w_2(A), w_2(B), w_1(B), w_3(B)$
- We can perform no swaps to this:
 - The first 2 operations are both on A and at least one is a write
 - The second and third operations are by the same transaction
 - The third and fourth are both on B at least one is a write and
 - So are the fourth and fifth
 - So this schedule is not conflict-equivalent to anything - and certainly not any serial schedules
- However, since nobody ever reads the values written by the $w_1(A), w_2(B)$ and $w_1(B)$ operations, the schedule has the same outcome as the serial outcome
 - $w_1(A), w_1(B), w_2(A), w_2(B), w_3(B)$



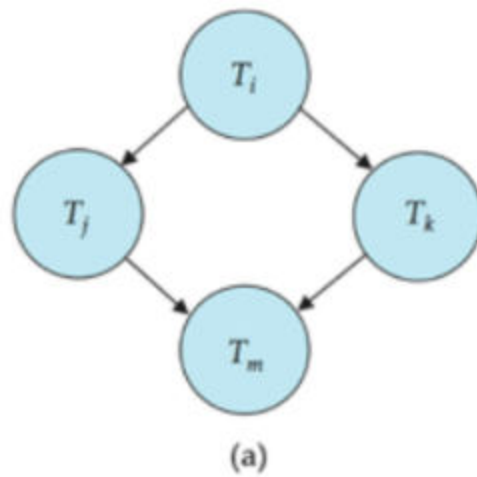
Precedence Graph

- Consider some schedule of a set of transactions T_1, T_2, \dots, T_n
- **Precedence Graph**
 - A directed graph where the vertices are the transactions (names)
- We draw an arc from T_i to T_j if the two transactions conflict and T_i accessed the data item on which the conflict arose earlier
- We may label the arc by the item that was accessed
- Example:



Testing for Conflict Serializability

- A schedule is conflict serializable if and only if its precedence graph is acyclic
- Cycle-detection algorithms exist which take order n^2 time, where n is the number of vertices in the graph
 - Better algorithms take order $n + e$ where e is the number of edges
- If precedence graph is acyclic, the serializability order can be obtained by a topological sorting of the graph
 - That is, linear order consistent with the partial order of the graph
 - For example, a serializability order for the schedule (a) would be one of either (b) or (c)




- Build a directed graph, with a vertex for each transaction
- Go through each operation of the schedule
 - If the operation is of the form $w_i(X)$, find each subsequent operation in the schedule also operating on the same data element X by a different transaction: that is, anything of the form $r_j(X)$ or $w_j(X)$
 - For each subsequent operation, add a directed edge in the graph from T_i to T_j
 - If the operation is of the form $r_i(X)$, find each subsequent write to the same data element X by a different transaction: that is, anything of the form $w_j(X)$
 - For each such subsequent write, add a directed edge in the graph from T_i to T_j
- The schedule is conflict-serializable if and only if the resulting directed graph is acyclic
- Moreover, we can perform a topological sort on the graph to discover the serial schedule to which the schedule is conflict-equivalent

-
- Consider the following schedule:
 - $w_1(A), r_2(A), w_1(B), w_3(C), r_2(C), r_4(B), w_2(D), w_4(E), r_5(D), w_5(E)$
 - We start with an empty graph with five vertices labeled T_1, T_2, T_3, T_4, T_5
 - We go through each operation in the schedule:

$w_1(A)$: A is subsequently read by T_2 , so add edge $T_1 \rightarrow T_2$
 $r_2(A)$: no subsequent writes to A , so no new edges
 $w_1(B)$: B is subsequently read by T_4 , so add edge $T_1 \rightarrow T_4$
 $w_3(C)$: C is subsequently read by T_2 , so add edge $T_3 \rightarrow T_2$
 $r_2(C)$: no subsequent writes to C , so no new edges
 $r_4(B)$: no subsequent writes to B , so no new edges
 $w_2(D)$: C is subsequently read by T_2 , so add edge $T_3 \rightarrow T_2$
 $w_4(E)$: E is subsequently written by T_5 , so add edge $T_4 \rightarrow T_5$
 $r_5(D)$: no subsequent writes to D , so no new edges
 $w_5(E)$: no subsequent operations on E , so no new edges

- We end up with a precedence graph



- This graph has no cycles, so the original schedule must be serializable
 - Moreover, since one way to topologically sort the graph is $T_3 - T_1 - T_4 - T_2 - T_5$, one serial schedule that is conflict-equivalent is

$$w_3(C), w_1(A), w_1(B), r_4(B), w_4(E), r_2(A), r_2(C), w_2(D), r_5(D), w_5(E)$$



Week 10 Lecture 3

▼ Class	BSCCS2001
🕒 Created	@November 8, 2021 5:01 PM
🔗 Materials	
≡ Module #	48
▼ Type	Lecture
# Week #	10

Transactions: Recoverability

What is Recovery?

- Serializability helps to ensure Isolation and Consistency of a schedule
- Yet, the Atomicity and Consistency may be compromised in the face of system failures
- Consider a schedule comprising of a single transaction (serial):
 - read(A)
 - $A := A - 50$
 - write(A)
 - read(B)
 - $B := B + 50$
 - write(B)
 - commit // Make the changes permanent; show the results to the user
- What if system fails after step 3 and before step 6?
 - Leads to inconsistent state
 - Need to rollback update of A
- This is known as Recovery

Recoverable Schedules

- If a transaction T_j reads a data item previously written by a transaction T_i , then the commit operation of T_i must appear before the commit operation of T_j
- The following schedule is not recoverable if T_9 commits immediately after the read(A) operation

T_8	T_9
read (A) write (A)	read (A) commit
read (B)	

- If T_8 should abort, T_9 would have read (and possibly shown to the user) an inconsistent database state
 - Hence, the database must ensure that schedules are recoverable

Cascading Rollbacks

- **Cascading rollback:** A single transaction failure leads to a series of transaction rollbacks
 - Consider the following schedule where none of the transactions has yet committed (so the schedule is recoverable)

T_{10}	T_{11}	T_{12}
read (A) read (B) write (A)	read (A) write (A)	read (A)
abort		

- If T_{10} fails, T_{11} and T_{12} must also be rolled back
- Can lead to the undoing of a significant amount of work

Cascadeless Schedules

- **Cascadeless schedules:** For each pair of transactions T_i and T_j such that T_j reads a data item previously written by T_i , the commit operation of T_i appears before the read operation of T_j
- Every cascadeless schedule is also recoverable
- It is desirable to restrict the schedules to those that are cascadeless
- Example of a schedule that is NOT cascadeless

T_{10}	T_{11}	T_{12}
read (A) read (B) write (A)	read (A) write (A)	read (A)
abort		

Example: Irrecoverable Schedule

T1	T1's Buffer	T2	T2's Buffer	Database
				A = 5000
R(A);	A = 5000			A = 5000
A = A – 1000;	A = 4000			A = 5000
W(A);	A = 4000			A = 4000
		R(A);	A = 4000	A = 4000
		A = A + 500;	A = 4500	A = 4000
		W(A);	A = 4500	A = 4500
		Commit;		
Failure Point				
Commit;				

Rollback is possible only till the end (commit) of T2
 So, the computation of A (4000) and write in T1 is lost

Example: Recoverable Schedule with Cascading Rollback

T1	T1's Buffer	T2	T2's Buffer	Database
				A = 5000
R(A);	A = 5000			A = 5000
A = A – 1000;	A = 4000			A = 5000
W(A);	A = 4000			A = 4000
		R(A);	A = 4000	A = 4000
		A = A + 500;	A = 4500	A = 4000
		W(A);	A = 4500	A = 4500
Failure Point				
Commit;				
		Commit;		

Rollback is possible as T2 has not committed yet
 But, T2 also need to be rolled back for rolling back T1

Example: Recoverable Schedule without Cascading Rollback

T1	T1's Buffer	T2	T2's Buffer	Database
				A = 5000
R(A);	A = 5000			A = 5000
A = A – 1000;	A = 4000			A = 5000
W(A);	A = 4000			A = 4000
Commit;				
		R(A);	A = 4000	A = 4000
		A = A + 500;	A = 4500	A = 4000
		W(A);	A = 4500	A = 4500
		Commit;		

Rollback is possible without cascading - wherever failure occurs

Transaction Definition in SQL

- Data manipulation language must include a construct for specifying the set of actions that comprise a transaction
 - In SQL, a transaction begins implicitly
 - A transaction in SQL ends by:
 - Commit work**
 - Commits current transaction and begins a new one
 - Rollback work**
 - Causes current transaction to abort
 - In almost al database systems, by default, every SQL statement also commits implicitly if it executes successfully
 - Implicit commit can be turned off by a database directive
 - For example in JDBC, `connection.setAutoCommit(false);`

Transaction Control Language (TCL)

- The following commands are used to control transactions
 - COMMIT**
 - To save the changes
 - ROLLBACK**
 - To roll back the changes
 - SAVEPOINT**
 - Creates points within the groups of transactions in which to ROLLBACK
 - SET TRANSACTION**
 - Places a name on a transaction
- Transactional control commands are only used with the DML Commands such as
 - INSERT, UPDATE and DELETE only
 - They cannot be used while creating tables or dropping them because these operations are automatically committed to the database

TCL: COMMIT Command

- COMMIT is the transactional command used to save changes invoked by a transaction to the database
- COMMIT saves all the transactions to the database since the last COMMIT or ROLLBACK command
- The syntax for the COMMIT command is as follows:

- `SQL> DELETE FROM Customers WHERE AGE = 25;`
- `SQL> COMMIT;`

`SQL> SELECT * FROM Customers;`

Before DELETE

ID	NAME	AGE	ADDRESS	SALARY
1	Ramesh	32	Ahmedabad	2000
2	Khilan	25	Delhi	1500
3	kaushik	23	Kota	2000
4	Chaitali	25	Mumbai	6500
5	Hardik	27	Bhopal	8500
6	Komal	22	MP	4500
7	Muffy	24	Indore	10000

`SQL> SELECT * FROM Customers;`

After DELETE

ID	NAME	AGE	ADDRESS	SALARY
1	Ramesh	32	Ahmedabad	2000
3	kaushik	23	Kota	2000
5	Hardik	27	Bhopal	8500
6	Komal	22	MP	4500
7	Muffy	24	Indore	10000

TCL: ROLLBACK Command

- The ROLLBACK is the command used to undo transactions that have not been already saved to the database
- This can only be used to undo transactions since the last COMMIT or ROLLBACK command was issued
- The syntax for a ROLLBACK command is as follows:
 - `SQL> DELETE FROM Customers WHERE AGE = 25;`
 - `SQL> ROLLBACK;`

`SQL> SELECT * FROM Customers;`

Before DELETE

ID	NAME	AGE	ADDRESS	SALARY
1	Ramesh	32	Ahmedabad	2000
2	Khilan	25	Delhi	1500
3	kaushik	23	Kota	2000
4	Chaitali	25	Mumbai	6500
5	Hardik	27	Bhopal	8500
6	Komal	22	MP	4500
7	Muffy	24	Indore	10000

`SQL> SELECT * FROM Customers;`

After DELETE

ID	NAME	AGE	ADDRESS	SALARY
1	Ramesh	32	Ahmedabad	2000
2	Khilan	25	Delhi	1500
3	kaushik	23	Kota	2000
4	Chaitali	25	Mumbai	6500
5	Hardik	27	Bhopal	8500
6	Komal	22	MP	4500
7	Muffy	24	Indore	10000

TCL: SAVEPOINT/ROLLBACK Command

- A SAVEPOINT is a point in a transaction when you can roll the transaction back to a certain point without rolling back the entire transaction
- The syntax for a SAVEPOINT command is
 - `SAVEPOINT SAVEPOINT_NAME;`
- This command serves only in the creation of a SAVEPOINT among all the transactional statements
- The ROLLBACK command is used to undo a group of transactions
- The syntax for rolling back to a SAVEPOINT is:
 - `ROLLBACK TO SAVEPOINT_NAME;`

Example:

- SQL> SAVEPOINT SP1;
 - Savepoint created.
- SQL> DELETE FROM Customers WHERE ID=1;
 - 1 row deleted.
- SQL> SAVEPOINT SP2;
 - Savepoint created.
- SQL> DELETE FROM Customers WHERE ID=2;
 - 1 row deleted.
- SQL> SAVEPOINT SP3;
 - Savepoint created.
- SQL> DELETE FROM Customers WHERE ID=3;
 - 1 row deleted.

- Three records deleted
- Undo the deletion of last two
- SQL> ROLLBACK TO SP2;
 - Rollback complete

```
SQL> SAVEPOINT SP1;
SQL> DELETE FROM Customers WHERE ID=1;
SQL> SAVEPOINT SP2;
SQL> DELETE FROM Customers WHERE ID=2;
SQL> SAVEPOINT SP3;
SQL> DELETE FROM Customers WHERE ID=3;
```

SQL> SELECT * FROM Customers

At the beginning

ID	NAME	AGE	ADDRESS	SALARY
1	Ramesh	32	Ahmedabad	2000
2	Khilan	25	Delhi	1500
3	kaushik	23	Kota	2000
4	Chaitali	25	Mumbai	6500
5	Hardik	27	Bhopal	8500
6	Komal	22	MP	4500
7	Muffy	24	Indore	10000

SQL> SELECT * FROM Customers;

After ROLLBACK

ID	NAME	AGE	ADDRESS	SALARY
2	Khilan	25	Delhi	1500
3	kaushik	23	Kota	2000
4	Chaitali	25	Mumbai	6500
5	Hardik	27	Bhopal	8500
6	Komal	22	MP	4500
7	Muffy	24	Indore	10000

TCL: RELEASE SAVEPOINT Command

- The RELEASE SAVEPOINT command is used to remove a SAVEPOINT that you have created
- The syntax for a RELEASE SAVEPOINT command is as follows
 - RELEASE SAVEPOINT SAVEPOINT_NAME;
- Once a SAVEPOINT has been released, you can no longer use the ROLLBACK command to undo transactions performed since the last SAVEPOINT

TCL: SET TRANSACTION Command

- The SET TRANSACTION command can be used to initiate a database transaction
- This command is used to specify a characteristics for the transactions that follows
 - For example, you can specify a transaction to be read-only or read-write
- The syntax for a SET TRANSACTION command is as follows:

- `SET TRANSACTION [READ WRITE | READ ONLY];`

View Serializability

- Let S and S' be two schedules with the same set of transactions
- S and S' are view equivalent if the following 3 conditions are met, for each data item Q
 - **Initial Read:** If in schedule S, transaction T_i reads the initial value of Q, then in schedule S' also transaction T_i must read the initial value of Q
 - **Write-Read Pair:** If in schedule S transaction T_i executed **read**(Q) and that value was produced by transaction T_j (if any), then in schedule S' also transaction T_i must read the value of Q that was produced by the same **write**(Q) operation of transaction T_j
 - **Final Write:** The transaction (if any) that performs the final **write**(Q) operation in schedule S must also perform the final **write**(Q) operation in schedule S'
- As can be seen, view equivalence is also based purely on reads and writes alone

- A schedule S is view serializable if it is view equivalent to a serial schedule
- Every conflict serializable schedule is also view serializable
- Below is a schedule which is view-serializable but not conflict serializable

T_{27}	T_{28}	T_{29}
read (Q)		
write (Q)	write (Q)	
		write (Q)

- What serial schedule is above equivalent to?
 - $T_{27} - T_{28} - T_{29}$
 - The one **read**(Q) instruction reads the initial value of Q in both schedules and
 - T_{29} performs the final write of Q in both schedules
- T_{28} and T_{29} perform **write**(Q) operations called blind writes, without having performed a **read**(Q) operation
- Every view serializable schedule that is not conflict serializable has blind writes

Test for View Serializability

- The %age graph test for conflict serializability cannot be used directly to test for view serializability
 - Extension to test for view serializability has cost exponential in the size of the precedence graph
- The problem of checking if a schedule is view serializable falls in the case of NP-complete problems
 - Thus, existence of an efficient algorithm is extremely unlikely
- However, practical assignments that just check some sufficient conditions for view serializability can still be used

View Serializability: Example 1

- Check whether the schedule is view serializable or not?
 - $S : R2(B); R2(A); R1(A); R3(A); W1(B); W2(B); W3(B)$
- Solution:
 - With 3 transactions, total number of schedules possible = $3! = 6$
 - $\langle T_1 T_2 T_3 \rangle$
 - $\langle T_1 T_3 T_2 \rangle$
 - $\langle T_2 T_3 T_1 \rangle$
 - $\langle T_2 T_1 T_3 \rangle$
 - $\langle T_3 T_1 T_2 \rangle$

- $\langle T_3 T_2 T_1 \rangle$

- Solution #2
 - Final update on data items:
 - A :- (No write on A)
 - B : T_1, T_2, T_3 (All 3 transactions write B)
 - As the final update on B is made by $T_3(T_1, T_2) \rightarrow T_3$
 - Now, removing those schedules in which T_3 is not executing at last:
 - $\langle T_1 T_2 T_3 \rangle$
 - $\langle T_2 T_1 T_3 \rangle$

- Solution #3
 - Initial Read + Which transaction updates after read?
 - A : T_2, T_1, T_3 (initial read)
 - B : T_2 (initial read); T_1 (update after read)
 - The transaction T_2 reads B initially which is updated by T_1
 - So, T_2 must execute before T_1
 - Hence, $T_2 \rightarrow T_1$
 - So, only one schedule survives:
 - $\langle T_2 T_1 T_3 \rangle$
 - Write Read Sequence (WR)
 - No need to check here
 - Hence, view equivalent serial schedule is:
 - $T_2 \rightarrow T_1 \rightarrow T_3$

View Serializability: Example 2

- Check whether S is Conflict serializable and / or view serializable or not?
 - $S : R1(A); R2(A); R3(A); R4(A); W1(B); W2(B); W3(B); W4(B)$

More Complex Notions of Serializability

- The schedule below produces the same outcome as the serial schedule $\langle T1, T5 \rangle$, yet is not conflict equivalent or view equivalent to it

T_1	T_5
read (A) $A := A - 50$ write (A)	
	read (B) $B := B - 10$ write (B)
read (B) $B := B + 50$ write (B)	
	read (A) $A := A + 10$ write (A)

- If we start with A = 1000 and B = 2000, the final result is 960 and 2040
- Determining such equivalence requires analysis of operations other than read and write



Week 10 Lecture 4

▼ Class	BSCCS2001
🕒 Created	@November 8, 2021 7:30 PM
🔗 Materials	
≡ Module #	49
▼ Type	Lecture
# Week #	10

Concurrency Control

- A database must provide a mechanism that will ensure that all possible schedules are both:
 - Conflict serializable
 - Recoverable and, preferably, Cascadeless
- A policy in which only one transaction can execute at a time generates serial schedules, but provides a poor degree of concurrency
- Concurrency-control schemes tradeoff between the amount of concurrency they allow and the amount of overhead that they incur
- Testing a schedule for serializability after it has executed is a little too late!
 - Tests for serializability help us understand why a concurrency control protocol is correct
- **Goal:** *To develop concurrency control protocols that will assure serializability*
- One way to ensure isolation is to require that data items be accessed in a mutually exclusive manner, that is, while one transaction is accessing a data item, no other transactions can modify that data item
 - Should a transaction hold a lock on the whole database
 - Would lead to strictly serial schedules - very poor performance
- The most common method used to implement locking requirement is to allow a transaction to access a data item only if it is currently holding a lock on that item


Lock-based Protocols

- A lock is a mechanism to control concurrent access to a data item


- Data items can be locked in two modes:
 - exclusive(X)* mode:
 - Data item can be both read as well as written
 - X-lock** is requested using **lock-X** instruction
 - shared(S)* mode:
 - Data item can only be read
 - S-lock** is requested using **lock-S** instruction
- A transaction can unlock a data item Q by the **unlock(Q)** instruction
- Lock requests are made to the concurrency-control manager by the programmer
- Transaction can proceed only after request is granted**

Lock-based Protocols: Lock Compatibility Matrix

- Lock-Compatibility Matrix: A lock compatibility matrix is used which states whether a data item can be locked by two transactions at the same time
- Full compatibility matrix

State of the lock	Lock request type 	
	Shared	Exclusive
Unlock	Yes	Yes
Shared	Yes	No
Exclusive	No	No

- Abbreviated compatibility matrix

State of the lock	Lock request type 	
	Shared	Exclusive
Shared	Yes	No
Exclusive	No	No

-
- Requesting for / Granting of a Lock
 - A transaction may be granted a lock on an item if the requested lock is compatible with locks already held on the item by other transactions
 - Sharing a Lock
 - Any number of transactions can hold shared locks on an item
 - But if any transaction holds an exclusive lock on the item no other transaction may hold any lock on the item
 - Waiting for a Lock
 - If a lock cannot be granted, the requesting transaction is made to wait till all incompatible locks held by other transactions have been released
 - Holding a Lock
 - A transaction must hold a lock on a data item as long as it accesses that item
 - Unlocking / Releasing a Lock
 - Transaction T_i may unlock a data item that it had locked at some earlier point
 - It is not necessarily desirable for a transaction to unlock a data item immediately after its final access of that data item, since serializability may not be ensured

Lock-Based Protocols: Example → Serial Schedule

- Let A and B be 2 accounts that are accessed by transactions T_1 and T_2
 - Transaction T_1 transfers \$50 from account B to account A
 - Transaction T_2 displays the total amount of money in accounts A and B, that is, the sum A + B
- Suppose that the values of accounts A and B are \$100 and \$200, respectively
- If these transactions are executed serially, either as T_1, T_2 or the order T_2, T_1 then transaction T_2 will display the value \$300

T1:

lock-X(B);
read(B);
B := B - 50;
write(B);
unlock(B);
lock-X(A);
read(A);
A := A + 50;
write(A);
unlock(A);

T2:

lock-S(A);
read(A);
unlock(A);
lock-S(B);
read(B);
unlock(B);
display(A + B)

Lock-Based Protocols: Example → Concurrent Schedule: Bad

- If, however, these transactions are executed concurrently, then schedule 1 is possible
- In this case, transaction T_2 displays \$250, which is incorrect
 - The reasons are ...
 - the transaction T_1 unlocked data item B too early, as a result of which T_2 saw an inconsistent state
 - Suppose we delay unlocking till the end

T1:

lock-X(B);
read(B);
B := B - 50;
write(B);
unlock(B);
lock-X(A);
read(A);
A := A + 50;
write(A);
unlock(A);

T2:

lock-S(A);
read(A);
unlock(A);
lock-S(B);
read(B);
unlock(B);
display(A + B)

T1	T2	Concurrency Control Manager
lock-X(B)		grant-X(B, T ₁)
read(B)		
B := B - 50		
write(B)		
unlock(B)		
	lock-S(A)	
	read(A)	grant-S(A, T ₂)
	unlock(A)	
	lock-S(B)	
	read(B)	grant-S(B, T ₂)
	unlock(B)	
	display(A + B)	
lock-X(A)		
read(A)		grant-X(A, T ₁)
A := A - 50		
write(A)		
unlock(A)		

Schedule 1

Lock-Based Protocols: Example → Concurrent Schedule: Good

- Delaying unlocking till the end, T_1 becomes T_3 & T_2 becomes T_4

T3:

lock-X(B);
read(B);
B := B - 50;
write(B);
lock-X(A);
read(A);
A := A + 50;
write(A);
unlock(B);
unlock(A)

T4:

lock-S(A);
read(A);
lock-S(B);
read(B);
display(A + B);
unlock(A);
unlock(B)


T ₁	T ₂	concurrency control manager
lock-X(B)		grant-X(B, T ₁)
read(B)		
B := B - 50		
write(B)		
unlock(B)		
	lock-S(A)	
	read(A)	grant-S(A, T ₂)
	unlock(A)	
	lock-S(B)	
	read(B)	grant-S(B, T ₂)
	unlock(B)	
	display(A + B)	
lock-X(A)		
read(A)		grant-X(A, T ₁)
A := A - 50		
write(A)		
unlock(A)		

Schedule 1

- Hence, sequence of reads and writes as in Schedule 1 is no longer possible
- T_4 will correctly display \$300

Lock-Based Protocols: Example → Concurrent Schedule: Deadlock

- Given T_3 and T_4 consider Schedule 2 (partial)
- Since T_3 is holding an exclusive mode lock on B and T_4 is requesting a shared-mode lock on B, T_4 is waiting for T_3 to unlock B
- Similarly, since T_4 is holding a shared-mode lock on A and T_3 is requesting an exclusive-mode lock on A, T_3 is waiting for T_4 to unlock A
- Thus, we have arrived at a state where neither of these transactions can ever proceed with its normal execution
- This situation is called a **deadlock**
- When deadlock occurs, the system must roll back one of the two transactions
- Once a transaction has been rolled back, the data items that were locked by that transaction are unlocked
- These data items are then available to the other transaction which can continue with its execution



Lock-Based Protocols

- If we do not use locking, or if we unlock data items too soon after reading or writing them, we may get inconsistent states
- On the other hand, if we do not unlock a data item before requesting a lock on another data item, deadlocks may occur
- Deadlocks are a necessary evil associated with locking, if we want to avoid inconsistent states
- Deadlocks are definitely preferable to inconsistent states, since they can be handled by rolling back transactions, whereas inconsistent states may lead to real-world problems that cannot be handled by the database system
- A **locking protocol** is a set of rules followed by all transactions while requesting and releasing locks
- Locking protocols restrict the set of possible schedules
- The set of all such schedules is a proper subset of all possible serializable schedules
- We present locking protocols that allow only conflict-serializable schedules, and thereby ensure isolation

Two-Phase Locking Protocol

- This protocol ensures conflict-serializable schedules
- Phase 1: Growing Phase
 - Transaction may obtain locks
 - Transaction may not release locks
- Phase 2: Shrinking Phase

- Transaction may release locks
- Transaction may not obtain locks
- The protocol assures serializability
 - It can be proved that the transactions can be serialized in the order of their lock points
 - That is, the point where a transaction acquires its final lock
- There can be conflict serializable schedules that cannot be obtained if two-phase locking is used
- However, in the absence of extra information (that is, ordering of access to data), two-phase locking is needed for conflict serializability in the following sense:
 - Given a transaction T_i that does not follow two-phase locking, we can find a transaction T_j that uses two-phase locking, and a schedule for T_i and T_j that is not conflict serializable

Lock Conversions

- Two-phase locking with lock conversions
 - First Phase
 - can acquire a lock-S on item
 - can acquire a lock-X on item
 - can convert a lock-S to a lock-X (upgrade)
 - Second Phase
 - can release a lock-S
 - can release a lock-X
 - can convert a lock-X to a lock-S (downgrade)
- This protocol assures serializability
 - But still relies on the programmer to insert the various locking instructions

Automatic Acquisition of Locks: Read

- A transaction T_i issues the standard read/write instruction, without explicit locking calls
- The operation **read**(D) is processed as:

```

if  $T_i$  has a lock on D
  then
    read(D)
  else begin
    if necessary, wait until no other transaction has a lock-X on D
    grant  $T_i$  a lock-S on D;
    read(D)
  end

```

Automatic Acquisition of Locks: Write

- **write**(D) is processed as:
- ```

if T_i has a lock-X on D
 then
 write(D)
 else begin
 if necessary, wait until no other transaction has any lock on D
 if T_i has a lock-S on D
 then

```

```
 upgrade lock on D to lock-X
 else
 grant T_i a lock-X on D
 write(D)

end;
```

- All locks are released after commit or abort

Deadlocks

- Two-phase locking does not ensure freedom from deadlocks

```
T3:
lock-X(B);
read(B);
B := B - 50;
write(B);
lock-X(A);
read(A);
A := A + 50;
write(A);
unlock(B);
unlock(A)
```

```
T4:
lock-S(A);
read(A);
lock-S(B);
read(B);
display(A + B);
unlock(A);
unlock(B)
```

| $T_3$       | $T_4$      |
|-------------|------------|
| lock-x (B)  |            |
| read (B)    |            |
| B := B - 50 |            |
| write (B)   |            |
|             | lock-s (A) |
|             | read (A)   |
|             | lock-s (B) |
| lock-x (A)  |            |

- Observe that transactions  $T_3$  and  $T_4$  are two phase, but, in deadlock

Starvation

- In addition to deadlocks, there is a possibility of **Starvation** (wot)
- **Starvation** occurs if the concurrency control manager is badly designed
  - For example:
    - A transaction may be waiting for an X-lock on an item, while a sequence of other transactions request and are granted an S-lock on the same item
    - The same transaction is repeatedly rolled back due to deadlocks
- Concurrency control manager can be designed to prevent starvation
- Starvation is also loosely referred to as **Livelock**

Cascading Rollback

- The potential for deadlock exists in most locking protocols
  - Deadlocks are necessary evil
- When a deadlock occurs there is a possibility of cascading roll-backs
- Cascading roll-back is possible under two-phase locking
- In the schedule here, each transaction observes the two-phase locking protocol, but the failure of T5 after the read(A) step of T7 leads to cascading rollback of T6 and T7



| $T_5$                                                                                         | $T_6$                                                         | $T_7$                        |
|-----------------------------------------------------------------------------------------------|---------------------------------------------------------------|------------------------------|
| lock-X( $A$ )<br>read( $A$ )<br>lock-S( $B$ )<br>read( $B$ )<br>write( $A$ )<br>unlock( $A$ ) | lock-X( $A$ )<br>read( $A$ )<br>write( $A$ )<br>unlock( $A$ ) | lock-S( $A$ )<br>read( $A$ ) |

### More Two Phase Locking Protocols

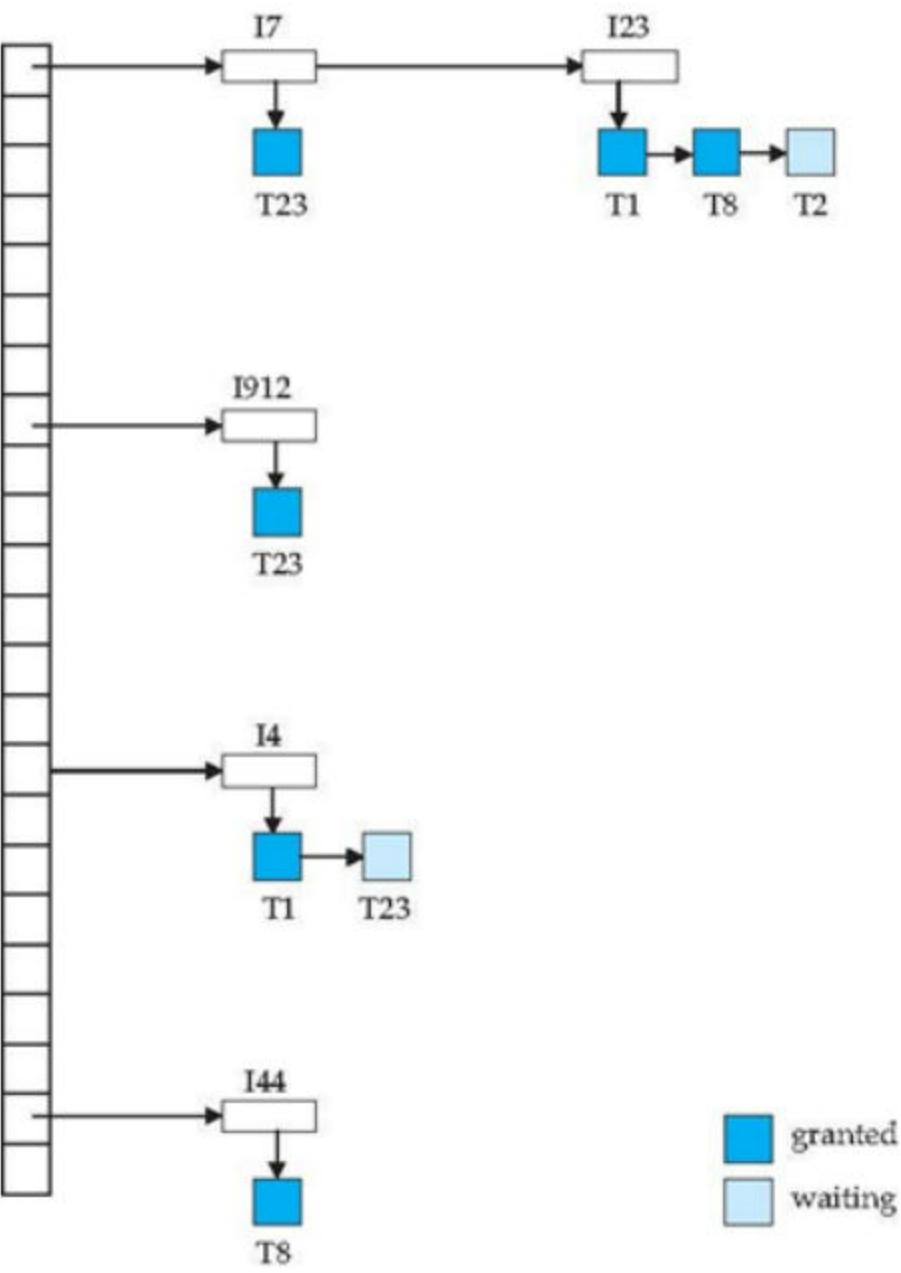
- To avoid Cascading roll-back, follow a modified protocol called strict two-phase locking
  - A transaction must hold all its exclusive locks till it commits/aborts
- Rigorous two-phase locking is even stricter
  - All locks are held till commit/abort
    - In this protocol, transactions can be serialized in the order in which they commit
- Note that concurrency goes down as we move to more and more strict locking protocol

### Implementation of Locking

- A lock manager can be implemented as a separate process to which transactions send lock and unlock requests
- The lock manager replies to a lock request by sending a lock grant messages (or a message asking the transaction to roll back, in case of a deadlock)
- The requesting transaction waits until its request is answered
- The lock manager maintains a data-structure called a lock table to record granted locks and pending requests
- The lock table is usually implemented as an in-memory hash table indexed on the name of the data item being locked

### Lock Table

- Dark blue rectangle indicate granted locks; light blue indicate waiting requests
- Lock table also records the type of lock granted or requested
- New request is added to the end of the queue of requests for the data item, and granted if it is compatible with all earlier locks
- Unlock requests result in the request being deleted, and later requests are checked to see it they can now be granted
- If transaction aborts, all waiting or granted requests of the transaction are deleted
  - Lock manager may keep a list of locks held by each transaction, to implement this efficiently





# Week 10 Lecture 5

|             |                           |
|-------------|---------------------------|
| ▼ Class     | BSCCS2001                 |
| 🕒 Created   | @November 9, 2021 4:41 PM |
| 🔗 Materials |                           |
| ☰ Module #  | 50                        |
| ▼ Type      | Lecture                   |
| # Week #    | 10                        |

## Concurrency Control (part 2)

### Deadlock Handling

- System is deadlocked if there is a set of transactions such that every transaction in the set is waiting for another transaction in the set
- Deadlock Prevention protocols ensure that the system will never enter into a deadlock state
  - Some prevention strats:
    - Require that each transaction locks all its data items before it beings execution (pre-declaration)
    - Impose partial ordering of all data items and require that a transaction can lock data items in the order specified by the partial order

### Deadlock Prevention

- **Transaction Timestamp:** Timestamp is a unique identifier created by the DBMS to identify the relative starting time of a transaction
  - Timestamping is a method of concurrency control in which each transaction is assigned a transaction timestamp
- Following schemes use transaction timestamps for the sake of deadlock prevention alone
  - **wait-die** scheme: non-preemptive
    - Older transaction may wait for younger one to release data item (here, older means smaller timestamp)
      - Younger transactions never wait for older ones; they are rolled back instead
    - A transaction may die several times before acquiring needed data item
  - **wound-wait scheme:** preemptive

- Older transaction wounds (forces rollback) of younger transaction instead of waiting for it
  - Younger transactions may wait for older ones
- May be fewer rollbacks than wait-die schemes

## Deadlock Prevention: Wait-Die Scheme

- It is a **non-preemptive** technique for deadlock prevention
- When transaction  $T_n$  requests a data item currently held by  $T_k$ ,  $T_n$  is allowed to wait only if it has a timestamp smaller than that of  $T_k$  (That is,  $T_n$  is older than  $T_k$ ), otherwise  $T_n$  is killed ("die")
- If a transaction requests to lock a resource (data item), which is already held with a conflicting lock by another transaction, then one of the two possibilities may occur:
  - **Timestamp( $T_n$ ) < Timestamp( $T_k$ ):**  $T_n$  which is requesting a conflicting lock, is older than  $T_k$ , then  $T_n$  is allowed to "wait" until the data-item is available
  - **Timestamp( $T_k$ ) > Timestamp( $T_n$ ):**  $T_n$  is younger than  $T_k$ , then  $T_n$  is killed ("dies")
    - $T_n$  is restarted later with a random delay but with the same timestamp( $n$ )
- This scheme allows the older transaction to "wait" but kills the younger one ("die")
- Example:
  - Suppose that transaction  $T_5, T_{10}, T_{15}$  have timestamps 5, 10 and 15 respectively
  - If  $T_5$  requests a data item held by  $T_{10}$  then  $T_5$  will "wait"
  - If  $T_{15}$  requests a data item held by  $T_{10}$ , then  $T_{15}$  will be killed ("die")

## Deadlock Prevention: Wound-Wait Scheme

- It is a preemptive technique for deadlock prevention
- When transaction  $T_n$  requests a data item currently held by  $T_k$ ,  $T_n$  is allowed to wait only if it has a timestamp larger than that of  $T_k$ , otherwise  $T_k$  is killed (wounded by  $T_n$ )
- If a transaction requests to lock a resource (data item), which is already held with a conflicting lock by another transaction, then one of the two possibilities may occur:
  - **Timestamp( $T_n$ ) < Timestamp( $T_k$ ):**  $T_n$  forces  $T_k$  to be killed ("wounds")
    - $T_k$  is restarted later with a random delay but with the same timestamp( $k$ )
    - **Timestamp( $T_n$ ) > Timestamp( $T_k$ ):**  $T_n$  "wait"s until the resource is free
- This scheme allows the younger transaction requesting a lock to "wait" if the older transaction already holds a lock, but forces the younger one to be suspended ("wound") if the older transaction requests a lock on an item already held by the younger one
- Example:
  - Suppose that transaction  $T_5, T_{10}, T_{15}$  have time-stamps 5, 10 and 15 respectively
  - If  $T_5$  requests a data item held by  $T_{10}$ , then it will be preempted from  $T_{10}$  and  $T_{10}$  will be suspended ("wounded")
  - If  $T_{15}$  requests a data item held by  $T_{10}$ , then  $T_{15}$  will "wait"



## Deadlock prevention

- Both in wait-die and in wound-wait schemes, a rolled back transaction is restarted with its original timestamp
  - Older transactions thus have precedence over newer ones, and starvation is hence avoided
- **Timeout-Based Schemes**
  - A transaction waits for a lock only for a specified amount of time
    - If the lock has not been granted within that time, the transaction is rolled back and restarted
  - Thus, deadlocks are not possible
  - Simple to implement; but starvation is possible
    - Also difficult to determine good value of the timeout interval

## Deadlock Detection

- Deadlocks can be described as a **wait-for graph**, which consists of a pair  $G = (V, E)$ 
  - $V$  is a set of vertices (all the transactions in the system)
  - $E$  is a set of edges; each element is an ordered pair  $T_i \rightarrow T_j$
- If  $T_i \rightarrow T_j$  is in  $E$ , then there is a directed edge from  $T_i$  to  $T_j$ , implying that  $T_i$  is waiting for  $T_j$  to release a data item
- When  $T_i$  requests a data item currently being held by  $T_j$ , then the edge  $T_i \rightarrow T_j$  is inserted in the wait-for graph
  - This edge is removed only when  $T_j$  is no longer holding a data item needed by  $T_i$
- The system is in a deadlock state if and only if the wait-for graph has a cycle
- Must invoke a deadlock-detection algorithm periodically to look for cycles

## Deadlock Detection: Example



## Deadlock Recovery

- When deadlock is detected:
  - Some transaction will have to be rolled back (made a victim) to break deadlock
    - Select that transaction as victim that will incur minimum cost
  - Rollback – determine how far to roll back transaction
    - Total rollback:** Abort the transaction and then restart it
    - More effective to roll back transaction only as far as necessary to break deadlock
  - Starvation happens if same transaction is always chosen as victim
    - Include the number of rollbacks in the cost factor to avoid starvation

## Timestamp-based Protocols

- Each transaction is issued a timestamp when it enters the system
  - If an old transaction  $T_i$  has time-stamp  $TS(T_i)$ , a new transaction  $T_j$  is assigned time-stamp  $TS(T_j)$  such that  $TS(T_i) < TS(T_j)$
- The protocol manages concurrent execution such that the time-stamps determine the serializability order
- In order to assure such behavior, the protocol maintains for each data  $Q$  two timestamp values:
  - W-timestamp(Q)** is the largest time-stamp of any transaction that executed write( $Q$ ) successfully
  - R-timestamp(Q)** is the largest time-stamp of any transaction that executed read( $Q$ ) successfully
- The timestamp ordering protocol ensures that any conflicting read and write operations are executed in timestamp order
- Suppose a transaction  $T_i$  issues a **read(Q)**
  - If  $TS(T_i) \leq \text{W-timestamp}(Q)$ , then  $T_i$  needs to read a value of  $Q$  that was already overwritten
    - Hence, the read operation is rejected, and  $T_i$  is rolled back



- If  $TS(T_i) \geq \mathbf{W}\text{-timestamp}(Q)$ , then the read operation is executed, and  $\mathbf{R}\text{-timestamp}(Q)$  is set to  $\max(\mathbf{R}\text{-timestamp}(Q), TS(T_i))$
- Suppose that transaction  $T_i$  issues **write**(Q)
  - If  $TS(T_i) < \mathbf{R}\text{-timestamp}(Q)$ , then the value of Q that  $T_i$  is producing was needed previously, and the system assumed that that value would never be produced
    - Hence, the **write** operation is rejected, and  $T_i$  is rolled back
  - If  $TS(T_i) < \mathbf{W}\text{-timestamp}(Q)$ , then  $T_i$  is attempting to write an obsolete value of Q
    - Hence, the **write** operation is rejected, and  $T_i$  is rolled back
  - Otherwise, the **write** operation is executed, and  $\mathbf{W}\text{-timestamp}(Q)$  is set to  $TS(T_i)$

## Example use of the Protocol

A partial schedule for several data items for transactions with timestamps 1, 2, 3, 4, 5

| $T_1$    | $T_2$             | $T_3$                  | $T_4$    | $T_5$                  |
|----------|-------------------|------------------------|----------|------------------------|
|          |                   |                        |          | read (X)               |
| read (Y) | read (Y)          |                        |          |                        |
|          |                   | write (Y)<br>write (Z) |          |                        |
|          | read (Z)<br>abort |                        |          | read (Z)               |
| read (X) |                   |                        | read (W) |                        |
|          |                   | write (W)<br>abort     |          |                        |
|          |                   |                        |          | write (Y)<br>write (Z) |

## Correctness of Timestamp-Ordering Protocol

- The timestamp-ordering protocol guarantees serializability since all the arcs in the precedence graph are of the form:



Thus, there will be no cycles in the precedence graph

- Timestamp protocol ensures freedom from deadlock as no transaction ever waits ~~(TATAKAE)~~
- But the schedule may not be cascade-free, may not even be recoverable